

# NYC\_taxi\_prediction

September 25, 2018

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
In [3]: data = pd.read_csv("./train.csv", sep=",")
```

```
In [4]: data.head()
```

```
Out[4]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21 UTC	
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16 UTC	
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00 UTC	
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42 UTC	
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00 UTC	

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	
2	-73.982738	40.761270	-73.991242	40.750562	
3	-73.987130	40.733143	-73.991567	40.758092	
4	-73.968095	40.768008	-73.956655	40.783762	

	passenger_count
0	1
1	1
2	2
3	1
4	1

```
In [5]: data.shape
```

```
Out[5]: (55423856, 8)
```

## 0.1 Data Cleaning Process

- Remove all the rows with NaN's.
- Remove all the rows with 0 values in each of the columns.
- Compute the 99.9 percentile and 0.1 percentile values for each of the latitudes and longitudes.
- Remove all the rows with values lying outside the above limits.
- Remove all the rows with fare amount  $\leq 0$  and also fare amount  $\geq 150$ USD (~99th percentile)
- Remove all the rows with passengers  $\geq 7$  (~99th percentile)

```
In [6]: missing_data = data[data["dropoff_longitude"].isnull() & data["dropoff_latitude"].isnull()]
```

```
In [7]: missing_data.shape
```

```
Out[7]: (376, 8)
```

```
In [8]: missing_data_inds = []  
        missing_data_inds += missing_data.index.tolist()
```

```
In [9]: zero_pickup_longitude = data[(data["pickup_longitude"]==0)]  
        zero_pickup_longitude.shape
```

```
Out[9]: (1055693, 8)
```

```
In [10]: missing_data_inds += zero_pickup_longitude.index.tolist()
```

```
In [11]: zero_pickup_latitude = data[(data["pickup_latitude"]==0)]  
        zero_pickup_latitude.shape
```

```
Out[11]: (1052158, 8)
```

```
In [12]: missing_data_inds += zero_pickup_latitude.index.tolist()
```

```
In [13]: zero_dropoff_longitude = data[(data["dropoff_longitude"]==0)]  
        zero_dropoff_longitude.shape
```

```
Out[13]: (1052745, 8)
```

```
In [14]: missing_data_inds += zero_dropoff_longitude.index.tolist()
```

```
In [15]: zero_dropoff_latitude = data[(data["dropoff_latitude"]==0)]  
        zero_dropoff_latitude.shape
```

```
Out[15]: (1049666, 8)
```

```
In [16]: missing_data_inds += zero_dropoff_latitude.index.tolist()
```

```
In [17]: zero_fare_amount = data[(data["fare_amount"]==0)]  
        zero_fare_amount.shape
```

```
Out[17]: (1380, 8)
```

```
In [18]: missing_data_inds += zero_fare_amount.index.tolist()
```

```
In [19]: zero_passenger_count = data[(data["passenger_count"]==0)]  
zero_passenger_count.shape
```

```
Out[19]: (195416, 8)
```

```
In [20]: missing_data_inds += zero_passenger_count.index.tolist()
```

```
In [21]: len(missing_data_inds)
```

```
Out[21]: 4407434
```

```
In [22]: missing_data_inds = list(set(missing_data_inds))
```

```
In [23]: len(missing_data_inds)
```

```
Out[23]: 1296797
```

```
In [24]: data_without_na_zero = data.drop(missing_data_inds)  
data_without_na_zero.shape
```

```
Out[24]: (54127059, 8)
```

```
In [25]: data_without_na_zero["fare_amount"].min()
```

```
Out[25]: -300.0
```

```
In [26]: data_without_na_zero["fare_amount"].max()
```

```
Out[26]: 93963.36
```

```
In [27]: data_without_na_zero = data_without_na_zero[data_without_na_zero["fare_amount"] >= 0]  
data_without_na_zero.shape
```

```
Out[27]: (54124853, 8)
```

```
In [28]: data_without_na_zero.describe()
```

```
Out[28]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
count	5.412485e+07	5.412485e+07	5.412485e+07	5.412485e+07	
mean	1.133736e+01	-7.391647e+01	4.069292e+01	-7.391484e+01	
std	2.085949e+01	7.863656e+00	7.508195e+00	7.890859e+00	
min	1.000000e-02	-3.442060e+03	-3.492264e+03	-3.442025e+03	
25%	6.000000e+00	-7.399228e+01	4.073652e+01	-7.399158e+01	
50%	8.500000e+00	-7.398209e+01	4.075335e+01	-7.398059e+01	
75%	1.250000e+01	-7.396826e+01	4.076754e+01	-7.396528e+01	
max	9.396336e+04	3.457626e+03	3.408790e+03	3.457622e+03	

	dropoff_latitude	passenger_count
count	5.412485e+07	5.412485e+07
mean	4.069177e+01	1.691539e+00
std	7.639962e+00	1.315083e+00
min	-3.493652e+03	1.000000e+00
25%	4.073551e+01	1.000000e+00
50%	4.075383e+01	1.000000e+00
75%	4.076839e+01	2.000000e+00
max	3.537133e+03	2.080000e+02

```
In [29]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["pickup_longitude"] > 0) & (data_without_na_zero["dropoff_longitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[29]: (54070729, 8)
```

```
In [30]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["pickup_longitude"] > 0) & (data_without_na_zero["dropoff_longitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[30]: (54033114, 8)
```

```
In [31]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["pickup_latitude"] > 0) & (data_without_na_zero["dropoff_latitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[31]: (53979080, 8)
```

```
In [32]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["pickup_latitude"] > 0) & (data_without_na_zero["dropoff_latitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[32]: (53925120, 8)
```

```
In [33]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["dropoff_longitude"] > 0) & (data_without_na_zero["dropoff_latitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[33]: (53871230, 8)
```

```
In [34]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["dropoff_longitude"] > 0) & (data_without_na_zero["dropoff_latitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[34]: (53817358, 8)
```

```
In [35]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["dropoff_latitude"] > 0) & (data_without_na_zero["dropoff_longitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[35]: (53763545, 8)
```

```
In [36]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["dropoff_latitude"] > 0) & (data_without_na_zero["dropoff_longitude"] > 0)) & (data_without_na_zero["passenger_count"] > 0)]
data_without_na_zero.shape
```

```
Out[36]: (53709782, 8)
```

```
In [37]: data_without_na_zero.describe()
```

```
Out [37]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
count	5.370978e+07	5.370978e+07	5.370978e+07	5.370978e+07	
mean	1.115886e+01	-7.397560e+01	4.075098e+01	-7.397462e+01	
std	2.065172e+01	3.376065e-02	2.631170e-02	3.230815e-02	
min	1.000000e-02	-7.404253e+01	4.063956e+01	-7.417725e+01	
25%	6.000000e+00	-7.399227e+01	4.073667e+01	-7.399156e+01	
50%	8.500000e+00	-7.398211e+01	4.075340e+01	-7.398063e+01	
75%	1.250000e+01	-7.396842e+01	4.076751e+01	-7.396558e+01	
max	9.396336e+04	-7.314683e+01	4.085129e+01	-7.374761e+01	

	dropoff_latitude	passenger_count
count	5.370978e+07	5.370978e+07
mean	4.075138e+01	1.691163e+00
std	2.965752e-02	1.313808e+00
min	4.060696e+01	1.000000e+00
25%	4.073585e+01	1.000000e+00
50%	4.075391e+01	1.000000e+00
75%	4.076836e+01	2.000000e+00
max	4.088127e+01	2.080000e+02

```
In [38]: data_without_na_zero = data_without_na_zero[((data_without_na_zero["passenger_count"] > 0) && data_without_na_zero.shape[0] > 0)]
```

```
Out [38]: (53709737, 8)
```

```
In [39]: data_without_na_zero["distance"] = np.sqrt(((data_without_na_zero["dropoff_latitude"] - data_without_na_zero["pickup_latitude"])**2 + (data_without_na_zero["dropoff_longitude"] - data_without_na_zero["pickup_longitude"])**2))
```

```
In [40]: data_without_na_zero.describe()
```

```
Out [40]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
count	5.370974e+07	5.370974e+07	5.370974e+07	5.370974e+07	
mean	1.115885e+01	-7.397560e+01	4.075098e+01	-7.397462e+01	
std	2.065172e+01	3.376064e-02	2.631170e-02	3.230814e-02	
min	1.000000e-02	-7.404253e+01	4.063956e+01	-7.417725e+01	
25%	6.000000e+00	-7.399227e+01	4.073667e+01	-7.399156e+01	
50%	8.500000e+00	-7.398211e+01	4.075340e+01	-7.398063e+01	
75%	1.250000e+01	-7.396842e+01	4.076751e+01	-7.396558e+01	
max	9.396336e+04	-7.314683e+01	4.085129e+01	-7.374761e+01	

	dropoff_latitude	passenger_count	distance
count	5.370974e+07	5.370974e+07	5.370974e+07
mean	4.075138e+01	1.691073e+00	3.342670e-02
std	2.965753e-02	1.307034e+00	3.717777e-02
min	4.060696e+01	1.000000e+00	0.000000e+00
25%	4.073585e+01	1.000000e+00	1.280780e-02
50%	4.075391e+01	1.000000e+00	2.173874e-02
75%	4.076836e+01	2.000000e+00	3.838472e-02
max	4.088127e+01	7.000000e+00	8.462889e-01

```

In [41]: data_without_na_zero = data_without_na_zero[data_without_na_zero["fare_amount"] <= 150]
         data_without_na_zero.shape

Out[41]: (53708886, 9)

In [42]: corr_distace_fare = data_without_na_zero["distance"].corr(data_without_na_zero["fare_

In [43]: corr_distace_fare

Out[43]: 0.883698656384901

In [44]: data_without_na_zero['pickup_datetime'] = data_without_na_zero['pickup_datetime'].str
         data_without_na_zero.shape

Out[44]: (53708886, 9)

In [45]: data_without_na_zero['pickup_datetime'] = pd.to_datetime(data_without_na_zero['pickup

In [46]: data_without_na_zero["time_of_day"] = (data_without_na_zero.pickup_datetime.dt.hour*60
         data_without_na_zero["time_of_day"].shape

Out[46]: (53708886,)

In [47]: corr_time_distance = data_without_na_zero["time_of_day"].corr(data_without_na_zero["d
         corr_time_distance


Out[47]: -0.029261452842436277

In [48]: corr_time_fare = data_without_na_zero["time_of_day"].corr(data_without_na_zero["fare_
         corr_time_fare


Out[48]: -0.01576489060431111

In [49]: plt.figure(figsize=(8,10))
         plt.xlabel("Time of the Day in minutes starting 00:00")
         plt.ylabel("Distance")
         plt.title("Scatter Plot Between Time of the day and Distance")
         plt.scatter(data_without_na_zero["time_of_day"][:1000], data_without_na_zero["distance
         plt.show()

```




```
In [50]: plt.figure(figsize=(8,10))
plt.xlabel("Distance")
plt.ylabel("Fare Amount")
plt.title("Scatter Plot Between Distance and Fare Amount")
plt.scatter(data_without_na_zero["distance"][:1000], data_without_na_zero["fare_amount"])
plt.show()
```



```
In [51]: plt.figure(figsize=(8,10))
plt.xlabel("Time of the Day in minutes starting 00:00")
plt.ylabel("Fare Amount")
plt.title("Scatter Plot Between Time of the day and Fare Amount")
plt.scatter(data_without_na_zero["time_of_day"][:1000], data_without_na_zero["fare_am
plt.show()
```





```
In [52]: data_without_na_zero["hour"] = data_without_na_zero.pickup_datetime.dt.hour
data_without_na_zero["day"] = data_without_na_zero.pickup_datetime.dt.day
data_without_na_zero["day_of_week"] = data_without_na_zero.pickup_datetime.dt.weekday
data_without_na_zero["month"] = data_without_na_zero.pickup_datetime.dt.month
```

```
In [53]: data_without_na_zero.head()
```

```
Out [53]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00	
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42	
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00	

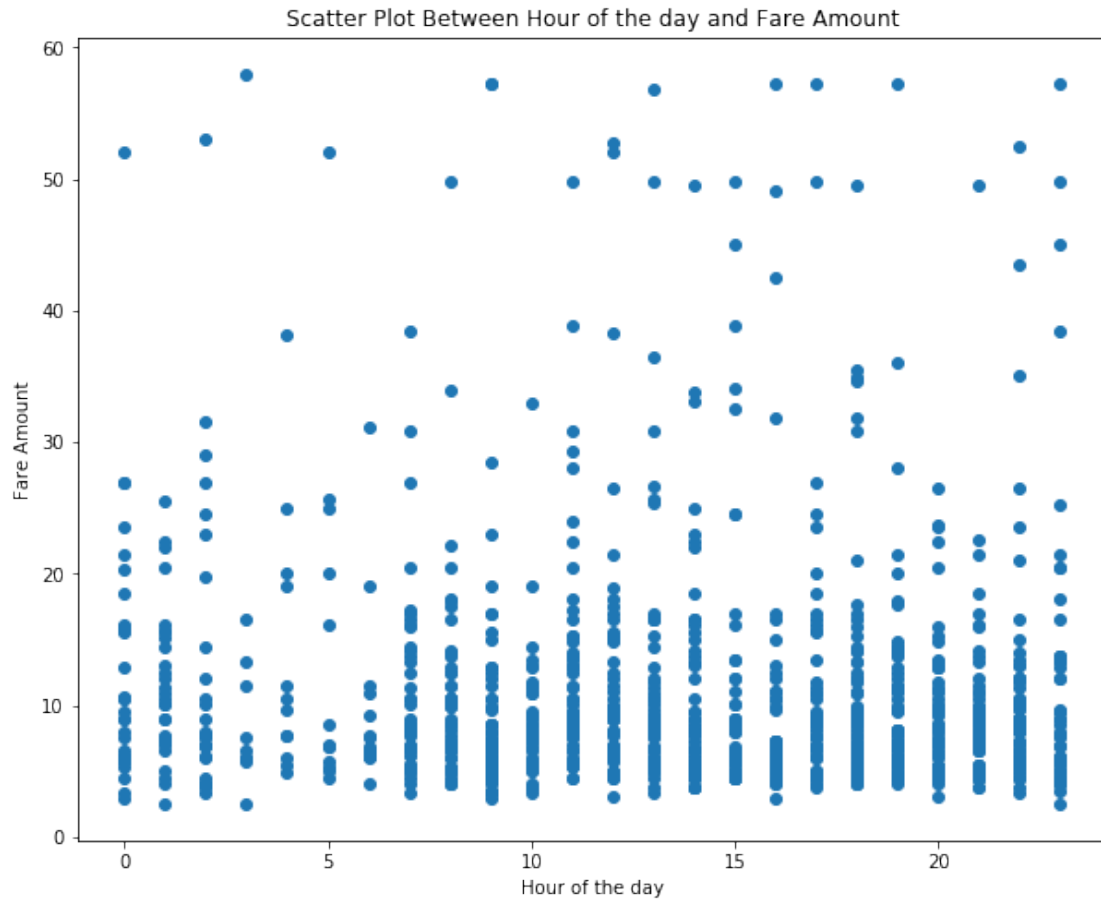
  

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	
2	-73.982738	40.761270	-73.991242	40.750562	
3	-73.987130	40.733143	-73.991567	40.758092	
4	-73.968095	40.768008	-73.956655	40.783762	

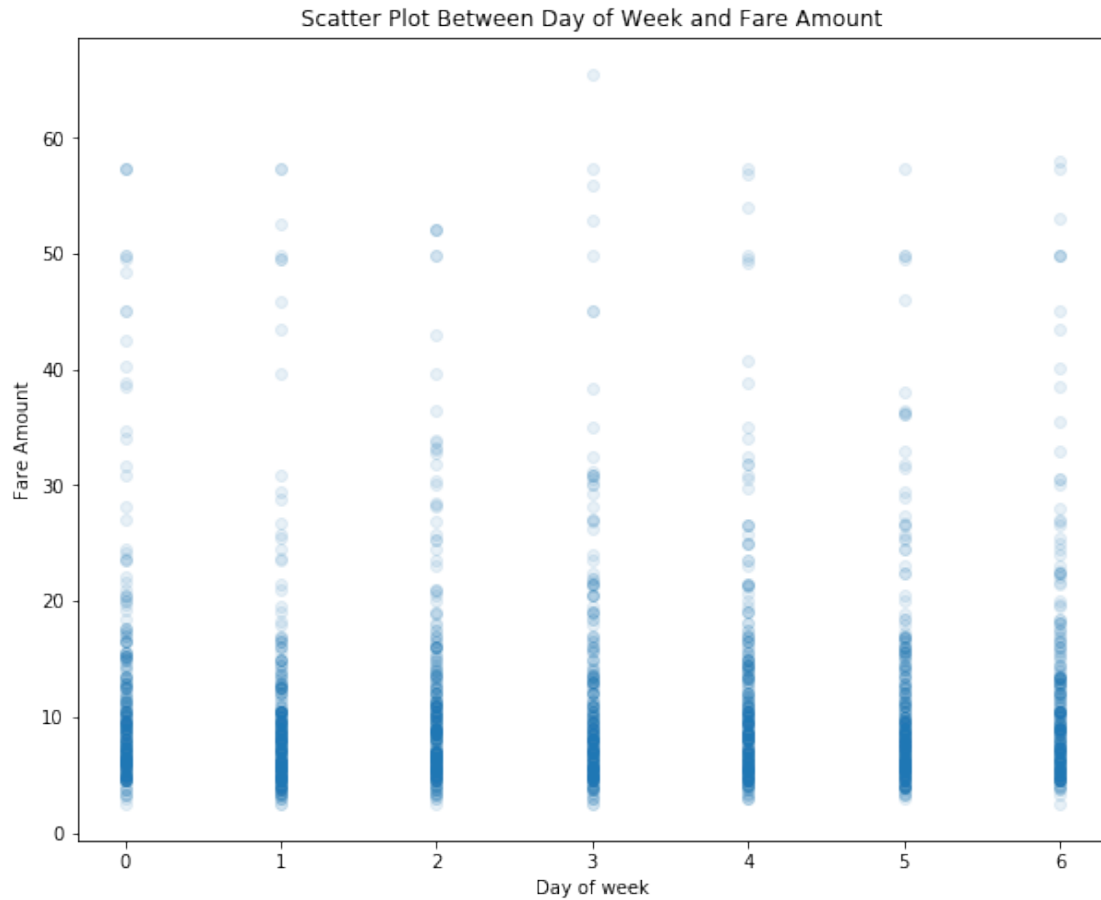
	passenger_count	distance	time_of_day	hour	day	day_of_week	month
0	1	0.009436	1046	17	15	0	6
1	1	0.079696	1012	16	5	1	1
2	2	0.013674	35	0	18	3	8
3	1	0.025340	270	4	21	5	4
4	1	0.019470	471	7	9	1	3

```
In [54]: plt.figure(figsize=(10,8))
plt.xlabel("Hour of the day")
plt.ylabel("Fare Amount")
plt.title("Scatter Plot Between Hour of the day and Fare Amount")
plt.scatter(data_without_na_zero["hour"][:1000], data_without_na_zero["fare_amount"])
plt.show()
```



- We observe that the prices of the cabs are high either during the late night and early mornings or during the peak office hours.

```
In [55]: plt.figure(figsize=(10,8))
plt.xlabel("Day of week")
plt.ylabel("Fare Amount")
plt.title("Scatter Plot Between Day of Week and Fare Amount")
plt.scatter(data_without_na_zero["day_of_week"][:2000], data_without_na_zero["fare_amount"][:2000])
plt.show()
```



```
In [56]: data_without_na_zero["day_of_week"] = data_without_na_zero.pickup_datetime.dt.day_name()
```

```
In [57]: data_without_na_zero.head()
```

```
Out[57]:
```

	key	fare_amount	pickup_datetime	\
0	2009-06-15 17:26:21.0000001	4.5	2009-06-15 17:26:21	
1	2010-01-05 16:52:16.0000002	16.9	2010-01-05 16:52:16	
2	2011-08-18 00:35:00.00000049	5.7	2011-08-18 00:35:00	
3	2012-04-21 04:30:42.0000001	7.7	2012-04-21 04:30:42	
4	2010-03-09 07:51:00.000000135	5.3	2010-03-09 07:51:00	


	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
0	-73.844311	40.721319	-73.841610	40.712278	
1	-74.016048	40.711303	-73.979268	40.782004	
2	-73.982738	40.761270	-73.991242	40.750562	
3	-73.987130	40.733143	-73.991567	40.758092	
4	-73.968095	40.768008	-73.956655	40.783762	

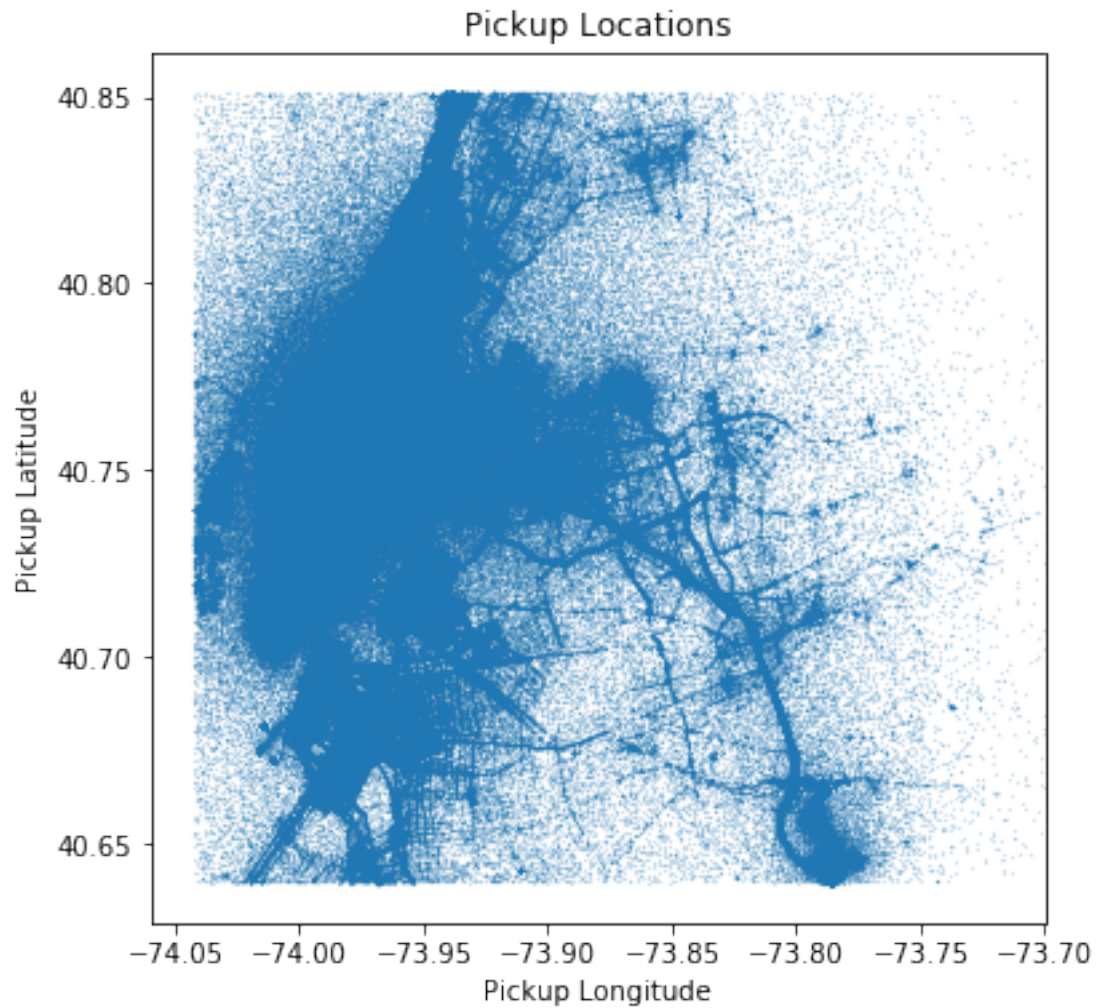
	passenger_count	distance	time_of_day	hour	day	day_of_week	month
--	-----------------	----------	-------------	------	-----	-------------	-------

0	1	0.009436	1046	17	15	Monday	6
1	1	0.079696	1012	16	5	Tuesday	1
2	2	0.013674	35	0	18	Thursday	8
3	1	0.025340	270	4	21	Saturday	4
4	1	0.019470	471	7	9	Tuesday	3

```
In [58]: plt.figure(figsize=(10,8))
plt.xlabel("Distance")
plt.ylabel("Day of week")
plt.title("Scatter Plot Between Distance and Day of Week")
plt.scatter(data_without_na_zero["distance"][:1000], data_without_na_zero["day_of_week"])
plt.show()
```




```
In [125]: plt.figure(figsize=(6,6))
plt.xlabel("Pickup Longitude")
plt.ylabel("Pickup Latitude")
plt.title("Pickup Locations")
plt.xlim(-74.06,-73.7)
plt.scatter(data_without_na_zero["pickup_longitude"], data_without_na_zero["pickup_latitude"])
plt.show()
```



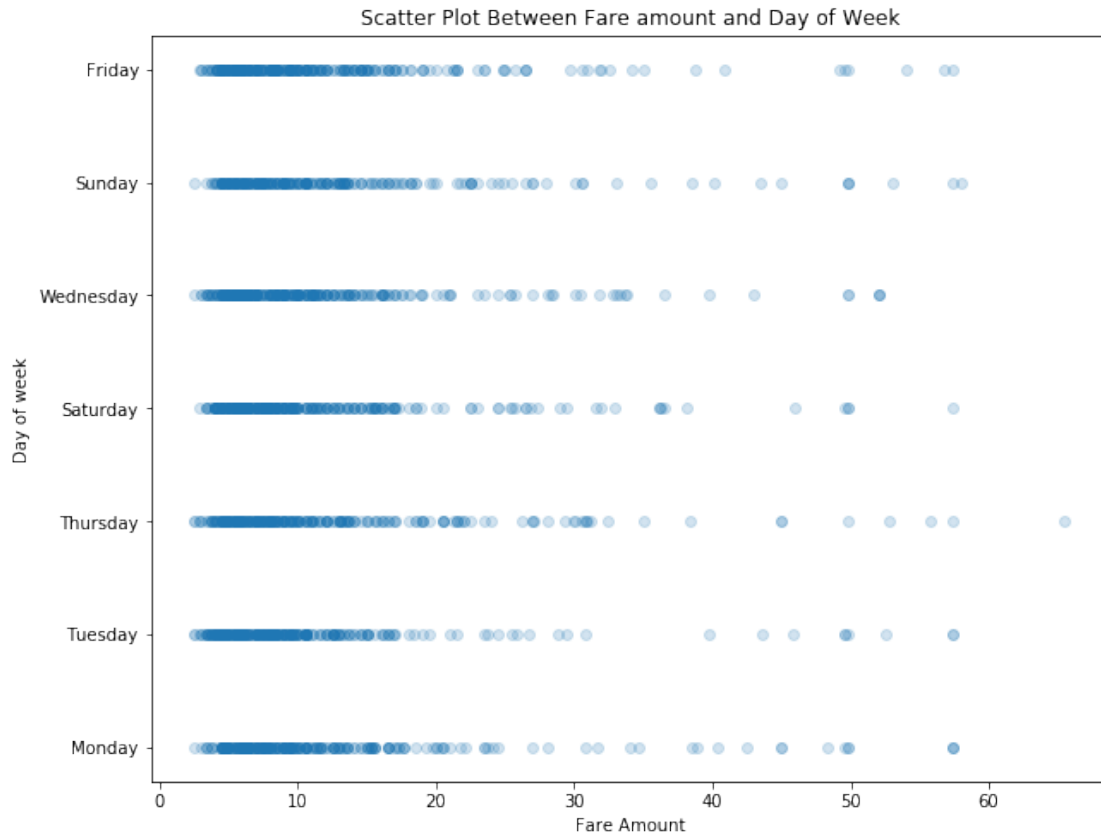
- We observe that most of the pickups are from Manhattan and Queens most of which includes airport pickups

```
In [124]: plt.figure(figsize=(6,6))
plt.xlabel("Dropoff Longitude")
plt.ylabel("Dropoff Latitude")
plt.title("Dropoff Locations")
plt.xlim(-74.06,-73.7)
plt.scatter(data_without_na_zero["dropoff_longitude"], data_without_na_zero["dropoff_latitude"])
plt.show()
```



- We observe that the drop off locations are distributed throughout all the Boroughs of NYC

```
In [61]: plt.figure(figsize=(10,8))
plt.xlabel("Fare Amount")
plt.ylabel("Day of week")
plt.title("Scatter Plot Between Fare amount and Day of Week")
plt.scatter(data_without_na_zero["fare_amount"][:2000], data_without_na_zero["day_of_week"])
plt.show()
```



```
In [62]: test_data = pd.read_csv("./test.csv", sep=",")
```

```
In [63]: test_data.head()
```

```
Out [63]:
```

	key	pickup_datetime	pickup_longitude \
0	2015-01-27 13:08:24.0000002	2015-01-27 13:08:24 UTC	-73.973320
1	2015-01-27 13:08:24.0000003	2015-01-27 13:08:24 UTC	-73.986862
2	2011-10-08 11:53:44.0000002	2011-10-08 11:53:44 UTC	-73.982524
3	2012-12-01 21:12:12.0000002	2012-12-01 21:12:12 UTC	-73.981160
4	2012-12-01 21:12:12.0000003	2012-12-01 21:12:12 UTC	-73.966046

	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	40.763805	-73.981430	40.743835	1
1	40.719383	-73.998886	40.739201	1
2	40.751260	-73.979654	40.746139	1
3	40.767807	-73.990448	40.751635	1
4	40.789775	-73.988565	40.744427	1

```
In [64]: data_without_na_zero["day_of_week"] = data_without_na_zero.pickup_datetime.dt.weekday
```

```
In [65]: data_without_na_zero.head()
```



```

Out [65]:
      key  fare_amount  pickup_datetime \
0  2009-06-15 17:26:21.0000001      4.5 2009-06-15 17:26:21
1  2010-01-05 16:52:16.0000002     16.9 2010-01-05 16:52:16
2  2011-08-18 00:35:00.00000049      5.7 2011-08-18 00:35:00
3  2012-04-21 04:30:42.0000001      7.7 2012-04-21 04:30:42
4  2010-03-09 07:51:00.000000135      5.3 2010-03-09 07:51:00

      pickup_longitude  pickup_latitude  dropoff_longitude  dropoff_latitude \
0      -73.844311      40.721319      -73.841610      40.712278
1      -74.016048      40.711303      -73.979268      40.782004
2      -73.982738      40.761270      -73.991242      40.750562
3      -73.987130      40.733143      -73.991567      40.758092
4      -73.968095      40.768008      -73.956655      40.783762

      passenger_count  distance  time_of_day  hour  day  day_of_week  month
0           1  0.009436      1046      17   15           0       6
1           1  0.079696      1012      16    5           1       1
2           2  0.013674       35      0   18           3       8
3           1  0.025340       270      4   21           5       4
4           1  0.019470       471      7    9           1       3

```

```
In [66]: # train_data = data_without_na_zero[:10000000]
```

```
In [67]: # target = train_data['fare_amount']
        # features = train_data.drop(['fare_amount', 'pickup_datetime', 'key'], axis=1)
```

```
In [68]: # X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=.2)
```

```
In [69]: # model = LinearRegression()
        # model.fit(features, target)
```

```
In [70]: # pred = model.predict(X_test)
```

- Below are the co-efficients for different features in the model

```
In [71]: # model.coef_
```

```
In [72]: # np.sqrt(metrics.mean_squared_error(y_test, pred))
```

```
In [73]: # X_train.head()
```

```
In [74]: test_data.describe()
```

```

Out [74]:
      pickup_longitude  pickup_latitude  dropoff_longitude  dropoff_latitude \
count      9914.000000      9914.000000      9914.000000      9914.000000
mean      -73.974722      40.751041      -73.973657      40.751743
std         0.042774      0.033541      0.039072      0.035435
min      -74.252193      40.573143      -74.263242      40.568973
25%      -73.992501      40.736125      -73.991247      40.735254
50%      -73.982326      40.753051      -73.980015      40.754065

```

75%	-73.968013	40.767113	-73.964059	40.768757
max	-72.986532	41.709555	-72.990963	41.696683

	passenger_count
count	9914.000000
mean	1.671273
std	1.278747
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	6.000000

```
In [75]: test_data["distance"] = np.sqrt(((test_data["dropoff_latitude"] - test_data["pickup_latitude"])
```

```
In [76]: test_data.head()
```

```
Out [76]:
```

	key	pickup_datetime	pickup_longitude	\
0	2015-01-27 13:08:24.0000002	2015-01-27 13:08:24 UTC	-73.973320	
1	2015-01-27 13:08:24.0000003	2015-01-27 13:08:24 UTC	-73.986862	
2	2011-10-08 11:53:44.0000002	2011-10-08 11:53:44 UTC	-73.982524	
3	2012-12-01 21:12:12.0000002	2012-12-01 21:12:12 UTC	-73.981160	
4	2012-12-01 21:12:12.0000003	2012-12-01 21:12:12 UTC	-73.966046	

	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	\
0	40.763805	-73.981430	40.743835	1	
1	40.719383	-73.998886	40.739201	1	
2	40.751260	-73.979654	40.746139	1	
3	40.767807	-73.990448	40.751635	1	
4	40.789775	-73.988565	40.744427	1	

	distance
0	0.021554
1	0.023180
2	0.005870
3	0.018649
4	0.050631

```
In [77]: test_data['pickup_datetime'] = test_data['pickup_datetime'].str.replace(" UTC", "")
test_data['pickup_datetime'] = pd.to_datetime(test_data['pickup_datetime'], format='%Y-%m-%d %H:%M:%S')
test_data["time_of_day"] = (test_data.pickup_datetime.dt.hour*60) + test_data.pickup_datetime.dt.minute
test_data["time_of_day"].shape
```

```
Out [77]: (9914,)
```

```
In [78]: test_data["hour"] = test_data.pickup_datetime.dt.hour
test_data["day"] = test_data.pickup_datetime.dt.day
test_data["day_of_week"] = test_data.pickup_datetime.dt.weekday
test_data["month"] = test_data.pickup_datetime.dt.month
```

```
In [79]: test_data.shape
```

```
Out[79]: (9914, 13)
```

```
In [80]: test_data.head()
```

```
Out[80]:
```

		key	pickup_datetime	pickup_longitude	\
0	2015-01-27 13:08:24.000000	2	2015-01-27 13:08:24	-73.973320	
1	2015-01-27 13:08:24.000000	3	2015-01-27 13:08:24	-73.986862	
2	2011-10-08 11:53:44.000000	2	2011-10-08 11:53:44	-73.982524	
3	2012-12-01 21:12:12.000000	2	2012-12-01 21:12:12	-73.981160	
4	2012-12-01 21:12:12.000000	3	2012-12-01 21:12:12	-73.966046	

	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	\
0	40.763805	-73.981430	40.743835	1	
1	40.719383	-73.998886	40.739201	1	
2	40.751260	-73.979654	40.746139	1	
3	40.767807	-73.990448	40.751635	1	
4	40.789775	-73.988565	40.744427	1	

	distance	time_of_day	hour	day	day_of_week	month
0	0.021554	788	13	27	1	1
1	0.023180	788	13	27	1	1
2	0.005870	713	11	8	5	10
3	0.018649	1272	21	1	5	12
4	0.050631	1272	21	1	5	12

```
In [81]: # x_test = test_data.drop(['pickup_datetime', 'key'], axis=1)
```

```
In [82]: # submission_pred = model.predict(x_test)
```

```
In [83]: # submission_pred
```

```
In [84]: # df = pd.DataFrame()
```

```
In [85]: # df["key"] = test_data["key"]
```

```
In [86]: # df["fare_amount"] = submission_pred
```

```
In [87]: # df.head()
```

```
In [88]: # df.to_csv("vr_linear_regression.csv", sep=",", index=False)
```

```
In [122]: train_data = data_without_na_zero[:10000000]
          target = train_data['fare_amount']
          features = train_data.drop(['fare_amount', 'pickup_datetime', 'key', 'time_of_day', 'day
          model = LinearRegression()
          model.fit(features, target)
          model.coef_
```

```
Out[122]: array([ 1.12453903e+01,  3.79198031e+00, -9.42194372e-01, -1.51826049e+01,
                  3.41778331e-02,  2.09508519e+02,  1.14840604e-02, -3.61266877e-02,
                  3.62673991e-02])
```

```
In [127]: features.columns
```

```
Out[127]: Index(['pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
                  'dropoff_latitude', 'passenger_count', 'distance', 'hour',
                  'day_of_week', 'month'],
                  dtype='object')
```

```
In [90]: x_test = test_data.drop(['pickup_datetime', 'key', 'time_of_day', 'day'], axis=1)
         submission_pred = model.predict(x_test)
         submission_pred
```

```
In [91]: df = pd.DataFrame()
         df["key"] = test_data["key"]
         df["fare_amount"] = submission_pred
         df.to_csv("vr_linear_regression.csv", sep=",", index=False)
```

```
In [92]: # from sklearn.ensemble import RandomForestRegressor
         # rf = RandomForestRegressor(n_estimators = 15)
         # rf.fit(features, target)
```

```
In [93]: # rf.feature_importances_
```

```
In [94]: # rf_5_preds = rf.predict(x_test)
```

```
In [95]: # df = pd.DataFrame()
         # df["key"] = test_data["key"]
         # df["fare_amount"] = rf_5_preds
         # df.to_csv("vr_Randomforest_regression.csv", sep=",", index=False)
```

```
In [96]: # from sklearn.svm import SVR
         # clf = SVR()
         # clf.fit(features, target)
```

```
In [97]: # from sklearn.neural_network import MLPRegressor
         # mlp = MLPRegressor(hidden_layer_sizes=(50,), activation='relu', solver='adam', learning_rate=0.001)
         # mlp.fit(features, target)
```

```
In [98]: # nn_preds = mlp.predict(x_test)
```

```
In [99]: # df = pd.DataFrame()
         # df["key"] = test_data["key"]
         # df["fare_amount"] = nn_preds
         # df.to_csv("vr_mlp_regressor.csv", sep=",", index=False)
```

```
In [100]: # target = (target - target.mean())/target.std()
```

```

In [101]: # features = (features - features.mean())/features.std()

In [102]: # rf = RandomForestRegressor(n_estimators = 10)
          # rf.fit(features, target)

In [103]: # x_test = (x_test - x_test.mean())/x_test.std()

In [104]: # x_test.head()

In [105]: # preds = rf.predict(x_test)

In [106]: # preds

In [107]: # preds = (preds*train_data["fare_amount"].std()) + train_data["fare_amount"].mean()

In [108]: # preds

In [109]: # df = pd.DataFrame()
          # df["key"] = test_data["key"]
          # df["fare_amount"] = preds
          # df.to_csv("vr_Randomforest_regression_v2.csv", sep=",", index=False)

In [115]: train_data = data_without_na_zero[:10000000]
          target = train_data['fare_amount']
          features = train_data.drop(['fare_amount', 'pickup_datetime', 'key', 'time_of_day', 'day']

In [116]: target = (target - target.mean())/target.std()
          features = (features - features.mean())/features.std()

In [117]: import lightgbm as lgbm
          lgbm_train_data = lgbm.Dataset(features, target, silent=True)
          params = {
              'boosting_type': 'gbdt', 'objective': 'regression', 'learning_rate': 0.005,
              'reg_alpha': 1, 'reg_lambda': 0.001, 'metric': 'rmse'}
          model = lgbm.train(params, train_set=lgbm_train_data, num_boost_round=1000)

In [120]: x_test = test_data.drop(['pickup_datetime', 'key', 'time_of_day', 'day'], axis=1)
          x_test = (x_test - x_test.mean())/x_test.std()
          lgbm_preds = model.predict(x_test)
          lgbm_preds = (lgbm_preds*train_data["fare_amount"].std()) + train_data["fare_amount"].
          print(lgbm_preds)

[ 8.99233897  9.34311555  5.3696885  ... 42.87808188 18.23889502
 6.92615164]

In [121]: df = pd.DataFrame()
          df["key"] = test_data["key"]
          df["fare_amount"] = lgbm_preds
          df.to_csv("vr_lgbm_v2.csv", sep=",", index=False)

```